

COMP90077 Assignment 2 Part I

(10% of the Subject Assessment)

An Experimental Study on Treaps

Due Date: **09:30AM May 23 (Monday), 2022**

Prepared by **Junhao Gan** and **Tony Wirth**

Background. When writing a research paper, sometimes, in addition to theoretical analysis, it is also important to include a section to demonstrate the *performance in practice* of the proposed algorithms (and data structures) by *experiments*. A good experiment section can often maximize the chance of a paper to be accepted, as it serves as extra convincing evidence for the superiority of the proposition in the paper.

In this Part of the assignment, you are asked to conduct an experimental study on Treaps and write an experiment section to demonstrate the results as if you are writing a research paper on this data structure.

Specifications. There are **three** main tasks in this assignment Part: (i) implement a data generator, a randomized treap and a simple competitor; (ii) conduct the required experiments; and (iii) write a report to demonstrate and analyse the experimental results.

Task 1: Implementations.

About Programming. There is *no* specified programming language in this assignment; you can use whatever programming language you prefer.

The Data Generator. In the experiment, each *data element* is a pair (id, key) , where *id* is a *unique identifier* and *key* is the *search key* of the element. Both *id* and *key* are **integers**. It may be helpful to keep an integer id_{next} to record the value of next identifier. Initially, $id_{next} \leftarrow 1$. More specifically, the data generator should have an interface, called *gen_element()*, to generate a new element by the following steps.

gen_element():

- $id \leftarrow id_{next}, id_{next} \leftarrow id_{next} + 1$;
- $key \leftarrow$ an integer that is drawn uniformly at random from the range $[0, 10^7]$;
- return element (id, key) ;

Since the identifiers of the elements are unique, when two elements happen to have a same search key, we can break the tie by taking the element with smaller *id* value as the smaller between the two. Furthermore, the data generator should also have three more interfaces, *gen_insertion()*, *gen_deleteion()* and *gen_search()*, to generate an insertion, a deletion and a search operation, respectively. Specifically, these operations have the following forms:

- an insertion of an element x : $(1, x)$, where 1 indicates that this operation is an insertion, and $x = (id, key)$;
- a deletion of a search key key_{del} : $(2, key_{del})$, where 2 indicates that this operation is a deletion, and key_{del} is the search key to be deleted; note that key_{del} is not necessarily in the current element set;

- a search of a key key_{sch} : $(3, key_{sch})$, where 3 indicates that this operation is a search, and key_{sch} is the search key to be searched.

The detailed steps of these three interfaces are as follows:

gen_insertion():

- $x \leftarrow$ a new element generated by invoking *gen_element()*;
- return $(1, x)$, an insertion for x ;

gen_deletion():

- $id_{del} \leftarrow$ an integer that is drawn uniformly at random from the range $[1, id_{next} - 1]$;
- if the element x with id_{del} has already been deleted,
 - $key_{del} \leftarrow$ an integer that is drawn uniformly at random from the range $[0, 10^7]$;
 - return $(2, key_{del})$, a deletion for deleting the search key key_{del} ;
- otherwise, return $(2, key_{del})$, a deletion for deleting the search key $x.key$ of the element x ;

As it is possible that more than one elements have the same search key, in a deletion, deleting an *arbitrary* element found with the target key is fine.

gen_search():

- $key_{sch} \leftarrow$ an integer that is drawn uniformly at random from the range $[0, 10^7]$;
- return $(3, key_{sch})$, a search operation with search key key_{sch} ;

The Randomized Treap. The randomized treap should be able to support three operations:

- *insert*(x): insert a new element x ;
- *delet*(key_{del}): delete an arbitrary element (if it exists) from the treap with search key key_{del} ;
- *search*(q): return an arbitrary element (if it exists) that has a search key $key = q$; otherwise, return NULL.

The Competitor. The competitor in this experiment is a *dynamic array*, where it supports the following operations:

- *insert*(x): call the *push-back* operation to insert element x ; note that the length will be doubled when the current array is full as discussed in class;
- *delete*(key_{del}): scan the array from the beginning; if an element x with search key key_{del} :
 - swap x with the *last* element at the back of the array;
 - delete x from the back;
 - if the current number of elements is smaller than $1/4$ length of the current array, *shrink* the length of the array into *half* by: (i) creating a new array with *half* length, (ii) copying all elements to the new array, and (iii) deleting the old array;

- $search(key_{sch})$: scan the array from the beginning and return the first element (if it exists) with search key $key = key_{sch}$; otherwise, return NULL.

Task 2: Experiments.

Conduct the following *four* experiments.

Exp 1: Time v.s. Number of Insertions. In this experiment, we study the *total running time* (of the randomized treap and the dynamic array, respectively) v.s. the lengths L_{ins} of *insertion-only* sequences. More specifically, we **first** generate *five* insertion-only sequences $\sigma_1, \sigma_2, \dots, \sigma_5$ respectively with $L_{ins} = 0.1M, 0.2M, 0.5M, 0.8M, 1M$, where M stands for *one million*, i.e., 10^6 . **Then** we input the **same** $\sigma_1, \dots, \sigma_5$ to each of the randomized treap and the dynamic array.

Exp 2: Time v.s. Deletion Percentage. In this experiment, we consider a sequence of a *fixed length* $L = 1M$ of updates (including both insertions and deletions). In particular, we vary the *in-expectation percentage*, $\%_{del}$, of the deletions in five update sequences $\sigma_1, \dots, \sigma_5$, where $\%_{del} = 0.1\%, 0.5\%, 1\%, 5\%, 10\%$. Specifically, an update sequence σ with $\%_{del}$ is generated as follows:

- $\sigma \leftarrow \emptyset$;
- for $i = 1$ to $L = 1M$,
 - generate an update s_i such that: s_i is a deletion with probability $\%_{del}$, and it is an insertion with probability $1 - \%_{del}$;
 - add s_i to σ ;
- return the update sequence σ ;

Run the randomized treap and the dynamic array on these sequences $\sigma_1 \dots, \sigma_5$, and measure their corresponding total running time.

Exp 3: Time v.s. Search Percentage. Analogous to Exp 2, we consider a sequence of a *fixed length* $L = 1M$ of operations which contains insertions and searches *only*. We vary the *rough percentage*, $\%_{sch}$, of the search operations in the operation sequence, where $\%_{sch} = 0.1\%, 0.5\%, 1\%, 5\%, 10\%$. An operation sequence σ with $\%_{sch}$ is generated in an analogous way as the update sequence with $\%_{del}$ in Exp 2. Run both of the two algorithms on these sequences and measure their corresponding overall running time on each of the sequences.

Exp 4: Time v.s. Length of Mixed Operation Sequence. In this experiment, we consider operation sequences that mix insertions, deletions and searches with different lengths $L = 0.1M, 0.2M, 0.5M, 0.8M, 1M$. Each of these sequences is generated such that each operation is a deletion with $\%_{del} = 5\%$, a search with $\%_{sch} = 5\%$, and an insertion with the rest probability. Measure the corresponding overall running time of both of the two algorithms on each of such sequences.

Task 3: The Report.

You will need to write and submit a report on the experimental results. More specifically, the report should contain the following components:

- Experiment Environment. Explain clearly the necessary information for others to *reproduce* your experiment, such as: the CPU frequency, memory, operating system, programming language, compiler (if applicable), and etc. **[0.5 mark out of 10]**

- Data Generation. Describe how you obtain or generate the data (i.e., operations in our context). Although in the above specification, certain details of the data generation are already given, you will need to describe how your data is generated with *your own words*. [0.5 mark out of 10]
- Experiments. [Each experiment takes 2 marks; thus in total 8 marks out of 10]
 - *Result Demonstration.* Report the results of the four experiments with diagrams. Specifically, in each diagram, the x -axis represents the variable: L_{ins} in Exp 1, $\%_{del}$ in Exp 2, $\%_{sch}$ in Exp 3, and L in Exp 4. The y -axis is the corresponding total running time. You will need to *plot* the results of the corresponding algorithms. Hence, there will be two lines in each diagram. [1 mark out of 2]
 - *Analysis.* In addition to showing the experimental results, you are also asked to *analyse* them. For example, you will need to explain, with certain theoretical analysis, why an algorithm performs better than the other in some cases while in other cases not. Sometimes, you will also need to analyse the results *across* different diagrams. [1 mark out of 2]
- Conclusion. Conclude your findings in these experiments. For example, in which cases, which algorithm is better. [1 mark out of 10]

Marking Scheme. The marks for each component in the report are as shown. The grading will be based on the *clarity* and *rigor* of your description or analysis for each part. This is individual work, and you must obey **academic integrity**. However, to ensure scientific integrity, indeed **reproducibility**, your descriptions of your experiments, their design and the design of the data structures must be sufficiently clear that a hypothetical algorithm developer could reproduce the whole experimental study with your report and obtain similar results.

Submissions. You should lodge your submission for Assignment 2 Part I via the LMS (i.e., Canvas). *You must identify yourself in **each** of your source files and the report.* Poor-quality scans of solutions written or printed on paper will *not* be accepted. There are scanning facilities on campus, not to mention scanning apps for smartphones etc. Solutions generated directly on a computer are of course acceptable. Submit *two* files:

- A *report.pdf* file, comprising your report for the experimental study.
- A *code.zip* file, containing all your sources files of the implementations for the experiments.

Do not include the testing files, as these might be large. **REPEAT: DO NOT INCLUDE TESTING FILES!** Moreover, it is very important, so that you can justify ownership of your work, that you detail your contributions in comments in your code, and in your report.

The Assignment continues on the next page!

COMP90077 Assignment 2 Part II (10% of the Subject Assessment)

Global Minimum Cuts

Due Date: **09:30AM May 23 (Monday), 2022**

In the week-8 lecture, and in the week-9 tutorial, we considered Karger's algorithm for Global Minimum Cuts. In this Part of the Assignment, you are to apply your competencies and skills to analyzing some details and extensions of Karger's Algorithm.

The questions

We have asked in class, but not that carefully, "What is the running time of the original algorithm?" We start by analyzing a baseline implementation of Karger's algorithm. As you show in part (a), the running time of one instance of the baseline version of Karger's algorithm is $O(n^2)$.

(a) [2 marks] There are $\Theta(n)$ contraction operations. **Show** that each contraction can be performed in $O(n)$ time. Consider a suitable simple data structure to store the graph edges.

(b) [2 marks] Suppose you wanted to speed up this algorithm, to run in much less than $\Theta(n^2)$ time. There is a well-studied data structure called *disjoint set* or *union find*. It has three operations:

Makeset Given an item x , $\text{Makeset}(x)$ produces a set with only the element x , its representative.

Find Given an item identity, x , $\text{Find}(x)$ returns the representative of the set that includes x .

Merge Given two item identities x and y , $\text{Merge}(x, y)$ replaces the two sets containing x and y with their union. The representative of the new set is the representative of one of the original sets: which depends on the implementation of Merge. If indeed x and y are in the same set, then there is no *effect*.

Observe that the Merge operation invokes the Find operation. To be sure, every item is an element of exactly one set, i.e., the sets are disjoint.

Describe how the union-find data structure is incorporated into the contraction operations of Karger's algorithm. *Hint:* For now, just keep a fixed array of the graph edges.

(c) [2 marks] One flaw with the setup above is that it might choose to contract an edge that is between two vertices that, due to other merges, are now in the same connected component. For example, suppose the graph originally has edges (x, y) , (y, z) , and (x, z) . Suppose (x, y) and (y, z) are contracted, then the edge (x, z) , if chosen for contraction, is in fact between two vertices already in the same "*super-vertex*". This "*false*" edge must therefore be skipped, and we choose a fresh candidate edge for contraction (which could also be false, in which case the process iterates).

Describe a simple technique to modify the array so that we do not waste too much time selecting edges between merged endpoints. *Hint:* This can be achieved by ensuring that once an edge is identified as *false*, it is excluded from future consideration, i.e., cannot again be a candidate for contraction.

(d) [2 marks] **Describe** the running time of this new union-find version of Karger's algorithm. How does it depend on the number of vertices, n , and the number of edges, m , and the running times of the Find and Merge operations, F and M , respectively?

(e) [1 mark, harder] There is another speedup for Karger’s algorithm. So far, we considered running $\Theta(n^2 \log n)$ copies of the contraction algorithm, and returning the smallest cut. This succeeds with probability $1/n^k$, for some k hidden in the Θ expression above.

Suppose we run the algorithm for $n - n/\sqrt{2}$ contractions. The probability that a particular global minimum cut remains in the graph is roughly

$$\frac{(n/\sqrt{2})^2}{n^2} = \frac{1}{2}.$$

Describe how many copies of the contraction algorithm would you run until $n/\sqrt{2}$ steps remain. Then how many until $n/2$ remain. Then how many until $n/(2\sqrt{2})$ remain. **Justify** your answer.

(f) [1 mark, even harder?] Based on your answer to (e), and that the running time of each step is the *basic* $O(n^2)$, **show** that the total asymptotic running time of this algorithm is $O(n^2 \log n)$.

Submissions. You should lodge your submission for Assignment 2 Part II via the LMS (i.e., Canvas). *You must identify yourself in your file.* Poor-quality scans of solutions written or printed on paper will *not* be accepted. There are scanning facilities on campus, not to mention scanning apps for smartphones etc. Solutions generated directly on a computer are of course acceptable. Submit *one* file:

- A *karger.pdf* file, comprising your solutions to the questions in this part.

Administrative Issues

When is late? What do I do if I am late? The due date and time are printed on the front of this document. The lateness policy is on the handout provided at the first lecture. As a reminder, the late penalty for non-exam assessment is *two* marks per day (or part thereof) overdue. Requests for extensions or adjustment must follow the University policy (the Melbourne School of Engineering “owns” this subject), including the requirement for appropriate evidence.

Late submissions should also be lodged via the LMS, but, as a courtesy, please also email both the two lecturers (Junhao Gan and Tony Wirth) when you submit late. If you make both on-time and late submissions, please consult the subject coordinators as soon as possible to determine which submission will be assessed.

Individual work. You are reminded that your submission for this Assignment is to be your own individual work. Students are expected to be familiar with and to observe the University’s Academic Integrity policy <http://academicintegrity.unimelb.edu.au/>. For the purpose of ensuring academic integrity, every submission attempt by a student may be inspected, regardless of the number of attempts made.

Students who allow other students access to their work run the risk of also being penalized, even if they themselves are sole authors of the submission in question. **By submitting your work electronically, you are declaring that this is your own work.** Automated similarity checking software may be used to compare submissions.

You may re-use code provided by the teaching staff, and you may refer to resources on the Web or in published or similar sources. Indeed, you are encouraged to read beyond the standard teaching materials. However, *all* such sources *must* be cited fully and, apart from code provided by the teaching staff, you must *not* copy code.

Finally. *Despite all these stern words, we are here to help!* There is information about getting help in this subject on the LMS pages. Frequently asked questions about the Assignment will be answered in the LMS discussion group.